

SKILL QUALITY REVIEW

purpose-check

Quality, efficacy, writing, and operational fit review

OVERALL VERDICT

85 / 100. Strong foundation with a valuable judgment model. Narrow the activation trigger and add deterministic decision gates before broad deployment.

Review item	Details
Artifact reviewed	purpose-check.skill.zip, inner payload purpose-check/SKILL.md
Payload size	7,456 UTF-8 bytes, 79 lines
Review scope	Writing quality, efficacy, trigger behavior, safety boundaries, packaging, and Hermes conventions
Review basis	Static source review plus predicted behavior across six representative scenarios

Prepared as a standalone review report

Executive summary

This is a strong, unusually thoughtful judgment skill. Its central flaw is serious but easy to fix: the frontmatter tells Hermes to invoke it for nearly every coding, writing, design, or generation request, while the body correctly argues for selective intervention and minimal friction. The body has the better judgment. The trigger needs to match it.

RECOMMENDATION

Keep the concept and most of the body. Narrow the trigger, add deterministic intervention gates and a verification checklist, and clarify the limits around autonomy and safety. Those changes should move this into the low 90s without sanding off its distinctive voice.

Scorecard

Area	Score	Assessment
Concept and purpose	18 / 20	Targets a real failure mode: producing the requested artifact while missing the user's actual objective.
Writing and organization	14 / 15	Memorable, polished, and coherent. The philosophy adds identity, with some cognitive load.
Efficacy and actionability	16 / 20	Four clear stages and strong failure modes. Decision gates and completion checks need sharpening.
Trigger precision and composability	10 / 15	The description invites near-universal activation and could compete with domain-specific skills.
Calibration and user experience	15 / 15	Excellent treatment of stakes, reversibility, supplied context, and minimal friction.
Metadata and conventions	6 / 10	Passes minimal validation, but misses documented peer conventions.
Safety and autonomy boundaries	6 / 10	Generally respectful, with some risk of paternalism around actual-good language.
Total	85 / 100	Strong foundation. Revise before broad deployment.

What the skill is trying to do

purpose-check introduces a judgment layer before and during build, generation, and creative work. It asks the agent to identify who benefits, test whether the literal request represents the real problem, define a verifiable end state, revisit important tradeoffs, and check whether the finished artifact still serves the original purpose.

The skill is built around four concepts: doing the requested work, doing it well, understanding the intended end, and applying practical judgment when context changes. Its strongest operating idea is that the agent should borrow missing context from the human early, before assumptions harden into implementation.

Strengths

1. It targets a real failure mode

Many agents optimize for literal compliance and artifact quality while missing whether the artifact should exist or whether it solves the right problem. This skill names that failure directly and gives the agent a repeatable way to catch it.

2. The four-stage structure is coherent

- Stage 1 asks whether the work should exist at all.
- Stage 2 defines what done means in terms a non-expert can verify.
- Stage 3 handles judgment at meaningful forks instead of treating defaults as destiny.
- Stage 4 checks for drift before delivery.

3. It handles friction better than most judgment prompts

The calibration section is excellent. It distinguishes a throwaway script from a relied-upon feature and from an action involving money, safety, other people's data, or irreversibility. The instruction to ask one good question instead of four ritual questions is exactly the right instinct.

4. It avoids exposing a philosophy checklist to the user

The skill explicitly warns the agent not to turn the conversation into a philosophy seminar. It recommends natural language such as “who ends up using this?” rather than internal terminology. That keeps the user experience grounded.

5. The failure modes are memorable and useful

- Placating: executing a request while suspecting it is the wrong request.
- Guessing the why: inventing a purpose to avoid asking.
- Ceremony creep: performing the stages as a ritual.
- End-loading judgment: noticing a fundamental mismatch only at delivery time.

Primary weakness: trigger precision

The description says to use the skill whenever the user asks the agent to code, create content, design a system, write a document, or generate any deliverable. That is effectively universal activation. The body then says that routine, reversible, well-specified tasks should receive only a ten-second internal pass.

This mismatch creates four risks: unnecessary questioning, latency on routine work, competition with specialized skills, and context pollution from a broad judgment layer that is present even when it has little to contribute. The description should route on ambiguity, consequence, or meaningful choice, not on the mere existence of a deliverable.

RECOMMENDED REPLACEMENT DESCRIPTION

Use when a request is ambiguous about its audience, intended outcome, constraints, or consequences; when materially different approaches are possible; or when new information suggests the requested artifact may not serve the stated goal. For routine, reversible, well-specified tasks, perform a brief internal check and proceed without questioning the user.

Recommended improvements

1. Add deterministic decision gates

The current skill relies heavily on judgment. That is appropriate for the subject, but a few explicit gates would make behavior more predictable without turning the skill into a rigid questionnaire.

Situation	Required behavior
Clear, low-risk, reversible work	Check internally and proceed.
Minor ambiguity with a safe default	State the material assumption and proceed.
Ambiguity that changes the deliverable	Ask one targeted question.
Money, privacy, safety, authorization, or irreversible action	Confirm purpose and constraints explicitly.
Likely request-to-goal mismatch	Flag it once, offer an alternative, then respect a permissible informed choice.

2. Add a verification hook

The skill is easier to admire than to evaluate because it does not define a compact completion check. Add an internal checklist that confirms:

- The intended audience or outcome is identified.
- Consequential assumptions are surfaced.
- The selected approach connects to the stated outcome.
- Unnecessary questions were avoided.
- Safety, privacy, authorization, and legal constraints still apply.
- Consequential mid-task tradeoffs were revisited before delivery.

This does not require exposing a decision log. A short assumption note is enough when visible confirmation would help.

3. Clarify autonomy and safety limits

The phrase “the actual good of the people involved” gives the skill moral weight, but it can also invite the model to substitute its own values for the user’s goals. Add a boundary stating that the skill aligns work with the user’s stated goals, constraints, and affected parties. It should not authorize moralizing or replacement of the user’s objective. It supplements safety, privacy, legal, and authorization checks, and it respects an informed choice only when the request remains permissible.

The existing instruction to say the quiet thing once and then respect the user’s call is good. The boundary should sit beside it.

4. Translate the philosophy into observable behavior

Keep the philosophical vocabulary because it gives the skill identity, but connect each term immediately to an action. Five Greek terms may be more than the model needs to operationalize in every invocation.

- Telos: identify the intended outcome.
- Phronesis: adapt the approach when context changes.
- Eudaimonia: consider affected people and foreseeable consequences.

This preserves the conceptual spine while reducing cognitive overhead.

5. Align with Hermes skill conventions

The official validator confirms that the file passes the hard requirements: frontmatter, name, description, description length, and non-empty body. The documented authoring convention is broader, however. Compared with established peer skills, this file should add:

- version
- author
- license
- platforms
- metadata.hermes.tags
- metadata.hermes.category
- metadata.hermes.related_skills

It should also consider the conventional structure of Overview, When to Use, actionable body, Common Pitfalls, and Verification Checklist. The current Failure modes to watch in yourself section is good content. Renaming or nesting it under the conventional heading would improve discoverability and consistency.

Writing quality review

What works stylistically

- The opening establishes a clear intellectual frame instead of defaulting to generic requirements language.
- The prose has a consistent voice and uses concrete examples to explain abstract ideas.
- The headings create a logical progression from purpose to execution to delivery.
- The skill repeatedly warns against ritual behavior, which is exactly the failure mode its own format could create.

Where the writing can tighten

- The frontmatter description is too long and repeats the body's operating model.
- The philosophical setup is memorable, but some of it could move into a short Overview so the operational rules arrive sooner.
- A few instructions are advisory rather than observable. Replace "answer these honestly" with explicit outputs or decision conditions where possible.
- The skill has no conventional When to Use or Verification Checklist heading, which makes it slightly harder to scan and compare with peers.
- The document would benefit from a short "Do not use this to" section covering routine, fully specified, low-risk tasks and requests already governed by a more specific safety or domain skill.

Safety and autonomy review

The skill is not an executable payload. The supplied archive contains a nested skill archive with one UTF-8 SKILL.md and benign macOS Finder metadata. The outer ZIP is unencrypted. The inner archive contains no executable files, symlinks, unsafe paths, or suspicious compression ratio.

The substantive risk is behavioral rather than code execution. Broad activation could make the agent over-question users or overrule clear requests. The skill already mitigates this with calibration and respect for informed choice. It needs a sharper statement that purpose alignment does not grant permission to make moral decisions on the user's behalf.

Static efficacy predictions

These are prompt-level predictions, not live behavioral test results.

Scenario	Predicted behavior
Sort records by date	Perform a brief internal check and proceed without asking a question.
Detailed product brief	Extract the supplied audience and constraints. Avoid re-asking for them.
Build a dashboard	Ask one high-value question about the audience or decision the dashboard supports.
Payment or customer-data workflow	Confirm purpose, authority, privacy constraints, and consequential assumptions.
Quick default versus better-fit design	Revisit the stated goal and surface the material tradeoff.
User says "use reasonable defaults"	Proceed for reversible, low-risk work. State material assumptions and pause for consequential ambiguity.

Packaging and metadata findings

- The outer ZIP is readable and unencrypted.
- It contains a nested purpose-check.skill archive plus __MACOSX Finder metadata.
- The inner archive contains purpose-check/SKILL.md only.
- The payload is valid UTF-8 and contains no executable code.
- The package is small, so context sprawl is not a concern at the file level.
- Removing Finder metadata would make distribution cleaner.
- The nested archive should be treated as an installer compatibility question only if the intended importer expects a direct SKILL.md directory or a different archive shape.

Prioritized action plan

1. Narrow the frontmatter description so the skill activates for ambiguity, consequence, or meaningful choice rather than every deliverable.
2. Add the five decision gates shown in this report.
3. Add a compact verification checklist and define what completion looks like for each stage.
4. Add an autonomy boundary that separates purpose alignment from moral substitution.
5. Add peer-matched metadata and conventional section headings.
6. Remove __MACOSX metadata from the distributable package.
7. Run a small behavioral evaluation set after revision, measuring unnecessary questions, missed clarifications, and correct handling of consequential ambiguity.

Final assessment

BOTTOM LINE

This is worth keeping. The idea is strong, the calibration is unusually good, and the writing has a point of view. The main defect is routing: the skill advertises itself as universal while its own body argues for selective use. Fix that mismatch, add explicit gates and a verification hook, and it becomes a much more dependable judgment layer.

Current score: 85 / 100. Expected score after the recommended revision: low 90s, assuming the revised trigger reduces unnecessary activation without increasing missed purpose checks.

Review method and confidence

The review was based on the actual nested SKILL.md payload, the archive structure, and Hermes skill-authoring conventions. The validator-level findings are confirmed from source. The efficacy section is a static prediction of likely behavior, so it should be followed by a small prompt evaluation after revision.

- Confirmed: the payload is valid UTF-8, non-executable, and free of unsafe archive paths or symlinks.
- Confirmed: the skill passes the documented minimum frontmatter requirements.
- High confidence: the activation description is broader than the body's intended calibration.
- Medium confidence: the proposed revision will reduce unnecessary questioning without weakening purpose checks.