SKILL QUALITY REVIEW

purpose-check

Fresh review of the enhanced skill

OVERALL VERDICT

96 / 100 Ready for broad deployment. The revision fixes every material weakness identified in the original review. A shorter description and one consolidation pass would make it exceptional.

Review item	Details
Artifact reviewed	purpose-check-enhanced.skill.zip, nested payload purpose-check/SKILL.md
Payload size	12,347 UTF-8 bytes, 132 lines
Review scope	Writing quality, efficacy, trigger behavior, safety boundaries, packaging, and Hermes conventions
Review basis	Static source review, mechanical validation, archive integrity checks, and scenario analysis
Review date	July 11, 2026

Prepared as a standalone review report

Executive summary

The enhanced version is a substantial improvement. It keeps the original skill's strongest idea, aligning literal work with the user's real objective, while turning that idea into predictable behavior. The revised trigger is selective. The intervention table tells the agent when to proceed, state an assumption, ask, confirm, or challenge. Safety and autonomy boundaries are explicit. The file now follows Hermes metadata and structural conventions.

RECOMMENDATION

Deploy it. The skill is safe, well-calibrated, and operationally useful. Before publishing a final release, shorten the frontmatter description and consolidate repeated lists of consequential domains. Neither issue blocks use.

Scorecard

Area	Score	Assessment
Concept and purpose	18 / 20	Still targets an important failure mode and now defines its boundary more clearly.
Writing and organization	13 / 15	Strong voice and scan-friendly structure. Some safety language repeats across sections.
Efficacy and actionability	18 / 20	Five observable gates and seven completion checks make the behavior testable.
Trigger precision and composability	14 / 15	Selective activation fixes the original routing flaw. The 469-character description remains heavier than necessary.
Calibration and user experience	15 / 15	Excellent handling of supplied context, reversibility, stakes, and one-question discipline.
Metadata and conventions	9 / 10	Complete peer-style metadata, conventional sections, valid YAML, and clean packaging.
Safety and autonomy boundaries	9 / 10	Clear limits on moralizing and informed choice. Repetition slightly obscures the governing rule.
Total	96 / 100	Production-ready with minor editorial cleanup remaining.

What changed since the first review

All seven recommended improvements were implemented: narrower routing, deterministic gates, a verification hook, explicit autonomy limits, observable translations of the philosophy, peer-matched metadata and structure, and removal of Finder metadata from the distributable package.

Resolution of the original findings

Original recommendation	Status	Evidence in enhanced skill
Narrow the activation trigger	**Resolved**	Frontmatter routes on material ambiguity, consequence, mismatch, and meaningful forks. Line 3.
Add deterministic decision gates	**Resolved**	Five situations map to observable behavior. Lines 39 to 49.
Add a verification hook	**Resolved**	Seven internal completion checks. Lines 122 to 132.
Clarify autonomy and safety limits	**Resolved**	The skill forbids moral substitution and preserves independent safety checks. Lines 21 to 25.
Translate philosophy into behavior	**Resolved**	Greek terms are immediately paired with actions and kept out of user-facing language. Lines 19 to 23.
Align with Hermes conventions	**Resolved**	Version, author, license, platforms, metadata, standard headings, pitfalls, and verification are present.
Clean the package	**Resolved**	The nested archive contains one SKILL.md and no Finder metadata.

Strongest improvements

1. The trigger now matches the body

The first review found a serious contradiction: universal activation in the description and selective intervention in the body. The new description routes on ambiguity, consequences, divergent approaches, goal mismatch, and consequential forks. Routine work gets a brief internal check without user friction.

2. Judgment is now observable

The intervention gate is the most important addition. It maps concrete situations to visible behavior and gives the model a clear stopping rule. This reduces the chance that the philosophy turns into ritual questioning.

3. The autonomy boundary is clear

The new text explicitly says that purpose alignment does not authorize moralizing or replacing the user's objective. It also states that informed choice applies only to permissible requests. That closes the largest behavioral safety gap in the original.

4. Completion can be evaluated

The verification checklist lets a reviewer test whether the skill identified the audience, outcome, assumptions, approach, consequential constraints, tradeoffs, and mismatches. The original had good ideas but no compact proof of completion.

Detailed quality review

Concept and purpose

The concept remains strong. Agents often satisfy a literal request while missing the decision, person, or outcome behind it. This skill introduces a compact judgment layer that catches that failure before implementation hardens around a bad assumption.

The enhanced version is also clearer about scope. It treats purpose checking as a routing and calibration mechanism rather than a competing domain process. That makes it more compatible with planning, engineering, writing, design, and safety skills.

Efficacy and actionability

The operational sequence is coherent:

1. Route through the intervention gate.
2. Decide whether the requested work should exist.
3. Define a verifiable end state.
4. Revisit purpose at consequential forks.
5. Check the finished artifact for drift.

Each stage now has a completion condition or a decision rule. The agent can tell when to continue without asking, when to state a reversible assumption, and when a response is required before proceeding.

> **WHY THIS SHOULD WORK**
>
> The skill ties intervention to material changes in the deliverable. That is the right unit of judgment. It avoids treating every missing detail as equally important.

Trigger precision and composability

The routing change deserves most of the score increase. The description now names the conditions that justify loading the skill and avoids the original invitation to activate on every deliverable.

One issue remains: the 469-character description combines activation criteria with two branches of operating behavior. That is valid and still under the 1,024-character limit, but it taxes every skill-selection pass. A tighter description could preserve routing while moving routine-default behavior into the body.

Suggested tighter description:

> Use when material ambiguity about audience, outcome, constraints, consequences, or approach could change the deliverable; when the requested output may not fit the underlying goal; or at consequential mid-task forks.

The body already handles routine checks and safe defaults well, so the selection description does not need to repeat them.

Writing and organization

The writing has a distinct point of view without becoming vague. The Greek vocabulary gives the skill a memorable internal model, and every term is translated into an action. The instruction to keep those terms out of user-facing conversation is smart and specific.

The new headings improve scanability. Overview, When to Use, Intervention Gate, Calibration, Common Pitfalls, and Verification Checklist match the way agents and reviewers look for operating rules.

Remaining editorial issue: repeated consequential lists

The categories money, privacy, safety, authorization, legal constraints, and irreversible actions recur throughout the description, overview, routing table, fork logic, calibration, pitfalls, and verification checklist. Repetition adds emphasis, but this amount creates maintenance risk. A future edit could change one list and leave the others stale.

Use one named definition near the top, such as **consequential domains**, then refer back to it. Keep the full list in the intervention gate and the verification checklist, where precision matters most.

Small points worth keeping

- "One question at a time" is a strong user-experience rule.
- "Rich context is an answer" directly prevents redundant clarification.
- The mismatch rule is calibrated well: flag once, explain the consequence, offer an alternative, then respect a permissible informed choice.
- The skill avoids visible internal logs unless the user asks for them.
- The `WHY.md` suggestion is limited to longer projects and explicitly excluded from throwaway work.

What should remain unchanged

Keep the intervention table, the one-question discipline, the distinction between safe defaults and consequential ambiguity, and the final drift check. Those are the parts most likely to change agent behavior in useful ways.

Safety, packaging, and validation

Behavioral safety

The skill now draws the right boundary around user autonomy. It aligns work with stated goals and affected parties while forbidding moral substitution. It preserves external safety, privacy, legal, and authorization checks. It also prevents a request to “use reasonable defaults” from bypassing consequential confirmation.

The remaining risk is over-application by a model that treats any design choice as material. The routing description and calibration section reduce that risk substantially. A small behavioral test suite should still measure unnecessary questions before release.

Package safety

The supplied outer ZIP is readable and unencrypted. It contains one nested `purpose-check.skill` archive. The nested archive contains one file, `purpose-check/SKILL.md`.

- No executable files
- No symlinks
- No absolute or traversal paths
- No URLs or outbound network references
- No shell commands or code fences
- No suspicious compression ratio
- No macOS Finder metadata
- ZIP integrity test passed

The payload SHA-256 is `0c81fbe9013d54689acac0b9696116bbcd6ab7988da7558a27ecb23f8fdb0ee9`.

Mechanical validation

Check	Result	Details
YAML frontmatter	**Pass**	Parses as a mapping and starts at byte zero.
Required fields	**Pass**	Name, description, version, author, license, platforms, and metadata present.
Description length	**Pass**	469 characters, below the 1,024-character limit.
Body and structure	**Pass**	Non-empty body, conventional headings, no duplicate level-two headings.
Verification coverage	**Pass**	Seven checkbox criteria.
Archive integrity	**Pass**	`unzip -t` reported no errors.
Encoding	**Pass**	Valid UTF-8.

Scenario analysis

These are static prompt-level predictions based on the enhanced instructions. They are stronger than the first review's predictions because the expected behavior now maps directly to explicit gate rows.

Scenario	Expected behavior
Sort records by date	Run a brief internal check and proceed without asking.
Detailed product brief	Extract the supplied audience, outcome, and constraints. Ask only if a missing distinction can change the deliverable.
Dashboard with no audience	Ask one targeted question about who uses it and which decision it supports.
Minor UI ambiguity	State the safe, reversible assumption briefly and proceed.
Payment or customer-data workflow	Confirm purpose, authority, affected parties, and relevant constraints before acting.
Requested artifact conflicts with stated goal	Flag the mismatch once, explain the likely consequence, and recommend a better fit.
User insists after hearing the trade-off	Respect the informed choice if the request remains permissible.
Mid-task privacy-changing fork	Pause and revisit the consequential tradeoff before implementation.

Recommended behavioral evaluation

Run a small benchmark before publishing the next version. Use at least 20 prompts split across four classes:

- clear and reversible tasks that should produce zero questions
- minor ambiguity with a safe default
- ambiguity that materially changes the deliverable
- consequential actions requiring explicit confirmation

Measure unnecessary-question rate, missed-clarification rate, redundant-question rate, and correct handling of informed choices. The most important regression test is whether routine tasks remain fast after the skill is installed alongside domain-specific skills.

Priority actions

1. Shorten the frontmatter description to activation criteria only.
2. Consolidate repeated consequential-domain lists into one definition plus references.
3. Run the behavioral benchmark and retain the prompts as regression fixtures.
4. Publish without further conceptual restructuring if the benchmark confirms low friction on routine work.

Final assessment

BOTTOM LINE

This revision is ready. It fixes the old routing contradiction, adds deterministic behavior, preserves user autonomy, and provides a real completion check. The remaining work is cleanup and measurement, not redesign.

The score rises from **85 / 100** to **96 / 100**. That increase is earned by concrete changes rather than polish alone. Every material recommendation from the first review appears in the enhanced source and can be tied to specific lines or package evidence.

The skill should be deployed broadly enough to test its value, with one guardrail: monitor whether it asks questions during clear, reversible work. If that rate stays low, the design is doing exactly what it should.

Review method and confidence

This review used the actual nested payload, archive structure, line-level source inspection, YAML parsing, static security checks, ZIP integrity verification, and comparison against Hermes skill-authoring conventions. No untrusted code was executed because the package contains no executable payload.

- **Confirmed:** valid UTF-8, safe archive paths, no symlinks, no executable content, and complete peer-style metadata.
- **Confirmed:** all seven recommendations from the original report are implemented.
- **High confidence:** trigger precision and user experience are materially improved.
- **High confidence:** safety and autonomy boundaries are adequate for deployment.
- **Medium confidence:** routine-task friction will remain low in live use. A behavioral benchmark is the right final proof.

Artifact identity

Outer ZIP SHA-256: `7233b20bb732db2d371d029f1e9aa35e47cf1a2dfab55b1cfb1dd6f6c95db92b`

Nested skill archive SHA-256: `ba293c98d386769b3c044b9a789ad622b66afbf0d2e0774c8a0e0c543528be7f`

SKILL.md SHA-256: `0c81fbe9013d54689acac0b9696116bbcd6ab7988da7558a27ecb23f8fdb0ee9`

Final verdict: deploy, then benchmark.