

purpose-check 1.2

Fresh structural and behavioral review

OVERALL VERDICT

89 / 100 Strong and close, but the internal stage contract is inconsistent. Fix the missing Stage 1 artifact, restore conventional metadata placement, and tighten the consequential-confirmation rule before broad deployment.

Review item	Details
Artifact reviewed	purpose-check-1.2 ZIP, payload purpose-check-1.2/SKILL.md
Payload size	9,859 UTF-8 bytes, 123 lines
Review scope	Structure, trigger behavior, internal dependencies, efficacy, safety, writing, packaging, and Hermes conventions
Review basis	Line-level source comparison, corpus convention check, archive inspection, minimum-validator reproduction, and scenario analysis
Review date	July 11, 2026

Prepared as a standalone review report

Executive summary

Version 1.2 is leaner, clearer, and better organized than 1.1. It defines consequential domains once, orders the intervention gate by consequence, removes a redundant calibration section, and cuts roughly 20 percent of the payload. Those are real improvements.

The revision also exposes a structural bug. Stage 3 tells the agent to check a choice against “your Stage 1 block,” and Stage 4 tells it to reread that block. The skill creates a purpose block only for longer projects and explicitly skips it for throwaway work. Medium-sized work can therefore enter Stages 3 and 4 without the named artifact those stages require.

Two smaller conflicts remain. Consequential work “always gets explicit confirmation,” while the skill also says rich supplied context is already an answer. The final delivery instruction requires one or two sentences of reasoning, which can add ceremony to clear routine work despite earlier instructions to proceed quietly.

RECOMMENDATION

Revise once more before broad deployment. The necessary changes are surgical: define a Stage 1 purpose anchor for every non-routine invocation, allow explicit supplied context to satisfy consequential confirmation where appropriate, make visible delivery reasoning conditional, and return version, author, and platforms to the top level of frontmatter.

Scorecard

Area	Score	Assessment
Concept and purpose	18 / 20	The judgment model remains valuable and distinctive.
Writing and organization	14 / 15	Sharper prose, less duplication, and better information hierarchy.
Efficacy and actionability	15 / 20	Clear gates, but later stages depend on an artifact that is not always created.
Trigger precision and composability	13 / 15	Selective routing is good. Consequential confirmation can still conflict with supplied-context reuse.
Calibration and user experience	13 / 15	Routine-task guidance is strong, but default delivery reasoning can reintroduce ceremony.
Metadata and conventions	7 / 10	Passes minimum validation. Version, author, and platforms are nested against dominant Hermes convention.
Safety and autonomy boundaries	9 / 10	Clear autonomy limits and consequential gates. Confirmation semantics need one clarification.
Total	89 / 100	Strong revision with one major structural defect and several minor conflicts.

What 1.2 improved

1. One definition replaces repeated domain lists

Lines 25 to 30 define **consequential domains** and **high-value question** once. Later sections reuse those terms. This fixes the repetition and drift risk present in 1.1. The definition is concrete enough to route behavior and short enough to remember.

2. Gate precedence is explicit

Lines 44 to 54 state that the table is ordered from lowest to highest consequence and that the highest applicable row wins. Version 1.1 said to apply the first row matching the highest consequence, which was harder to parse. The new rule is clearer.

3. The payload is tighter

The source shrank from 12,347 bytes and 132 lines to 9,859 bytes and 123 lines. The deleted Calibration section mostly repeated the gate and definitions. Removing it improves information hierarchy without losing the central behavior.

4. The writing sounds more natural

The prose is less academic and more direct. It drops the ergon and arete labels, keeps the three contextual terms that carry useful distinctions, and improves several examples. “Concrete enough to disagree with” is stronger than generic advice to make a purpose note testable.

5. The gate and checklist remain usable

The five gate rows still distinguish quiet execution, reversible assumptions, deliverable-changing ambiguity, likely goal mismatch, and consequential action. The six verification checks are compact and mostly observable.

BOTTOM LINE ON THE REVISION

Fable improved the skill’s density and local clarity. The score is held back by cross-section contracts, where instructions depend on terms or artifacts established only under narrower conditions.

Comparison with 1.1

Area	1.1	1.2
Consequential vocabulary	Repeated lists	Single named definition
Gate precedence	Awkward first-row wording	Explicit low-to-high ordering
Payload	12,347 bytes	9,859 bytes
Stage artifact	Purpose block only for long work	Same condition, now referenced as mandatory later
Metadata	Conventional top-level fields	Fields moved under generic metadata

Primary structural conflict

Undefined Stage 1 block

The skill creates a **purpose block** only for longer projects:

Line 71: “For longer projects, write a short purpose block in your working notes or a `WHY.md` ... Skip this for throwaway tasks.”

Later stages treat a **Stage 1 block** as an existing artifact:

Line 88: “Check the choice against your Stage 1 block.”

Line 98: “Reread your Stage 1 block.”

These terms are close enough that a model may infer the intended connection, but the creation rule and consumption rule do not match. The gap is largest for medium-sized work: a feature, document, or dashboard substantial enough to reach Stage 3, yet short enough that the agent reasonably skipped a written purpose block.

Why it matters

The skill’s core mechanism is preserving purpose through forks and final review. If the purpose anchor does not exist, Stage 3 and Stage 4 fall back to memory or improvisation. That weakens the exact behavior the skill was designed to make reliable.

This is a structural defect because one stage consumes an artifact under broader conditions than the stage that produces it. It is not fixed by better prose elsewhere.

Recommended fix

Create a purpose anchor for every invocation that passes beyond the routine internal check. The anchor can live in working memory and does not need a file:

Suggested contract: “Before Stage 2, hold a one-sentence purpose anchor: who this is for and what gets better. For longer projects, write it in working notes or `WHY.md`. For shorter work, keep it internal. Stages 3 and 4 refer back to this same anchor.”

Then replace both “Stage 1 block” references with “purpose anchor.” One term, one creation rule, one lifecycle.

Severity

Major. The defect does not make the skill unsafe, but it breaks the internal dependency behind its central promise. Fixing it should recover three to four points by itself.

Secondary conflicts

Consequential confirmation versus rich context

Line 27 says work in consequential domains **always gets explicit confirmation**. Line 28 says rich context is an answer and supplied constraints should be extracted instead of asked again. Line 54 again requires explicit confirmation.

These rules can conflict when a detailed request already states the purpose, affected parties, authority, and constraints. A literal reading forces a redundant question. A context-first reading permits proceeding. The skill needs to choose.

Recommended rule: supplied context may satisfy confirmation when it is explicit, current, and attributable to an authorized user. Ask only for missing or uncertain consequential facts. Keep a separate confirmation requirement for actions that independently require fresh consent under the governing safety or domain policy.

Delivery reasoning versus quiet routine work

Lines 42 and 50 tell the agent to perform a brief internal check and proceed without ceremony. Line 103 says, “Then deliver, with one or two sentences of the reasoning that mattered.” That instruction has no condition, so it can make every result carry an explanation.

Change it to: “At delivery, include one or two sentences about a material purpose or judgment call only when they help the user evaluate the result.”

Verification lost the fork check

Version 1.1 explicitly required consequential mid-task tradeoffs to be revisited. Version 1.2 keeps Stage 3 but removes that check from the final checklist. Since consequential forks are a frontmatter trigger, this omission weakens completion coverage.

Add one checkbox: “Material mid-task forks were checked against the purpose anchor and surfaced when they changed the deliverable or consequential exposure.”

Informed choice lost its safety qualifier in one pitfall

Line 107 says to flag the wrong request once, offer the alternative, and “respect their call.” The intervention gate and verification checklist correctly limit this to permissible requests. The pitfall should use the same qualifier to avoid a local contradiction.

EDITORIAL PRINCIPLE

Each governing rule should have one canonical statement. Short references can point back to it. Rewriting the same rule with slightly different qualifiers is how safety and autonomy language drifts.

Metadata, packaging, and mechanical validation

Frontmatter convention

The file passes the documented minimum requirements: byte-zero frontmatter, valid YAML mapping, name, description under 1,024 characters, and a non-empty body. The description is 386 characters.

Version 1.2 moves `version`, `author`, and `platforms` beneath generic metadata while leaving `license` at the top level. This is valid YAML, but it breaks the established peer shape and creates an odd split among package identity fields.

A corpus scan of 224 installed skills found:

Field placement	Top level	Nested
<code>version</code>	121	2
<code>author</code>	96	2
<code>platforms</code>	100	0

Restore all three fields beside name, description, and license. The `metadata.hermes` namespace should continue to hold tags, category, and related skills.

Package safety

The attachment has no filename extension but is a valid ZIP. It contains one regular file, `purpose-check-1.2/SKILL.md`.

- Archive is unencrypted
- ZIP integrity test passed
- No executable files
- No symlinks
- No absolute or traversal paths
- No URLs or outbound network references
- No shell commands
- No suspicious compression ratio

Ad-hoc validation

A focused temporary validator reproduced Hermes's documented minimum checks and was removed after execution.

Check	Result	Details
UTF-8 and size	Pass	9,859 bytes, below 100,000-byte limit.
Frontmatter	Pass	Valid YAML mapping with required name and description.
Description	Pass	386 characters, below 1,024-character limit.
Body	Pass	Non-empty, six level-two sections, no duplicate headings.
Archive integrity	Pass	<code>unzip -t</code> reported no errors.
Peer metadata shape	Warn	Version, author, and platforms are nested.

Scenario analysis

These are deterministic readings of the written rules, not live model evaluation results.

Scenario	Expected route	Conflict
Clear sort function	Internal check, proceed without asking.	None. Lines 42 and 50 agree.
Minor reversible UI default	State one-line assumption and proceed.	None. Line 51 is clear.
Dashboard missing audience	Ask one high-value question.	None. Lines 28 and 52 agree.
Payment workflow with a complete brief	Confirm consequential facts.	“Always confirm” can require a redundant question despite the rich-context rule.
Medium feature with a meaningful fork	Check the fork against Stage 1 and bring a recommendation.	The named Stage 1 block may not exist.
Routine task delivery	Quiet completion.	Line 103 can require visible reasoning anyway.

Recommended behavioral benchmark

The next revision should be tested with at least 24 prompts across these classes:

- six clear, reversible tasks
- six minor ambiguities with safe defaults
- six deliverable-changing ambiguities
- six consequential tasks, half with complete authorization context and half with missing context

Measure unnecessary questions, missed clarifications, redundant confirmations, visible ceremony on routine work, correct purpose-anchor creation, and correct handling of informed permissible choices.

The critical regression fixture is a medium-sized task with a material mid-task fork. The agent should create or retain a purpose anchor before the fork, compare the options against it, and surface only the decision that changes the user-visible result.

Prioritized revision plan

1. Define one purpose anchor and establish when it is created.
2. Replace both Stage 1 block references with that term.
3. Clarify when supplied context satisfies consequential confirmation.
4. Make delivery reasoning conditional on usefulness.
5. Restore the mid-task fork checkbox.
6. Restore version, author, and platforms to top-level frontmatter.
7. Qualify every “respect their call” instruction with permissibility.

Final assessment

BOTTOM LINE

Version 1.2 is worth keeping and refining. Its prose and information hierarchy are better than 1.1. The central stage dependency is still broken, and two local rules can reintroduce redundant questions or visible ceremony. Fix those contracts before calling it production-ready.

The fresh score is **89 / 100**. This aligns with Fable's headline score because the structural findings independently reproduce from the source. The agreement is incidental; the evidence is what matters.

My prior 96 for version 1.1 was too high. I confirmed that recommended sections existed but did not trace their producer-consumer contracts closely enough. In particular, I missed that the later stages referenced a Stage 1 artifact that was only conditionally created. This review corrects that mistake and weights cross-section coherence more heavily.

Review method and confidence

The review used the actual ZIP payload, a line-level comparison against 1.1, archive traversal and symlink checks, YAML parsing, a 224-skill metadata corpus scan, a temporary minimum-validator reproduction, and a six-scenario rule matrix. No untrusted code was executed because the package contains only Markdown.

- **Confirmed:** valid ZIP, safe paths, one UTF-8 Markdown payload, and passing minimum frontmatter validation.
- **Confirmed:** the Stage 1 block is referenced on lines 88 and 98 but created only for longer projects on line 71 under a different name.
- **Confirmed:** metadata placement diverges sharply from the installed Hermes corpus.
- **High confidence:** the consequential-confirmation wording can cause redundant questions.
- **Medium confidence:** live models will expose these conflicts consistently. A behavioral benchmark is still needed.

Artifact identity

Attachment SHA-256: 8da8cd7ba30a6dd672825a0232373a5fa2317aa38fb2a6cb9798183dddb4f337

SKILL.md SHA-256: da13319b03b98ce3b0d347d63eb15d416b2e439f8e3d8140da9c556f82cc2f07

Final verdict: revise the contracts, then benchmark.