

purpose-check 1.3

Fresh behavioral and cross-platform packaging review

OVERALL VERDICT

95 / 100 Behaviorally production-ready. Version 1.3 resolves the structural defects from 1.2. The supplied artifact passes Anthropic’s packager and Hermes’s minimum loader. A separate build variant is needed for strict open-spec portability.

Review item	Details
Artifact reviewed	purpose-check-1.3 ZIP, payload purpose-check-1.3/SKILL.md
Payload size	10,702 UTF-8 bytes, 124 lines
Review scope	Behavior, internal contracts, safety, writing, Anthropic packaging, Hermes loading, and Agent Skills reference portability
Review basis	Line-level comparison, authoritative validator source, three-validator testing, archive inspection, and scenario analysis
Review date	July 11, 2026

Prepared as a standalone review report

Executive summary

Version 1.3 closes the behavioral defects identified in 1.2. It defines a purpose anchor for every invocation that advances beyond the routine internal check, uses that same term in Stages 3 and 4, lets explicit current context satisfy consequential confirmation, restores the mid-task fork verification, and makes delivery reasoning conditional.

Fable’s pushback on the prior review was correct. Anthropic’s packager allows only name, description, license, allowed-tools, metadata, and compatibility at the top level. Moving version, author, and platforms to the top level would fail that validator. The Agent Skills specification itself demonstrates author and version under metadata.

Fable was also right that a skill cannot authenticate an “authorized user.” The 1.3 wording uses facts the agent can evaluate: supplied context may satisfy confirmation when it is explicit and current; the agent asks only for missing, stale, or uncertain consequential facts.

The remaining issue is distribution portability. The current artifact passes Anthropic’s validator and Hermes’s documented minimum. It fails the official skills-ref reference validator because metadata contains flow-style lists and nested structured values. Its versioned directory name also differs from name: purpose-check, which the open specification rejects.

RECOMMENDATION

Approve 1.3 for Anthropic and Hermes use. Keep this source as the canonical behavioral document. Add a small build step that emits target-specific frontmatter and packages the directory as purpose-check/ for strict Agent Skills consumers.

Scorecard

Area	Score	Assessment
Concept and purpose	19 / 20	Distinctive judgment model with a clear purpose anchor.
Writing and organization	14 / 15	Lean, direct, and well structured. One long definition carries heavy policy density.
Efficacy and actionability	19 / 20	Producer-consumer contracts now line up across all four stages.
Trigger precision and composability	14 / 15	Selective routing and context reuse are coherent. Live trigger-rate evidence is still absent.
Calibration and user experience	15 / 15	Routine work stays quiet; consequential gaps get focused questions.
Packaging and conventions	6 / 10	Passes Anthropic and Hermes. Current artifact fails strict open-spec validation.
Safety and autonomy boundaries	8 / 10	Strong canonical boundaries. Live behavioral testing remains the final proof.
Total	95 / 100	Behaviorally ready with a targeted portability build task remaining.

Resolution of the 1.2 findings

1.2 finding	Status	Evidence in 1.3
Undefined Stage 1 block	Resolved	Line 71 creates one purpose anchor for every non-routine invocation. Lines 88, 98, and 122 consume the same term.
Confirmation conflicts with supplied context	Resolved	Lines 27 and 54 state that explicit supplied context counts as confirmation.
Delivery reasoning adds ceremony	Resolved	Line 103 makes visible reasoning conditional and says routine work ships quietly.
Fork verification missing	Resolved	Line 122 restores the material mid-task fork check.
Top-level metadata recommendation	Withdrawn	Anthropic rejects unexpected top-level fields. Nested author and version match the Agent Skills spec.
Repeat permissibility in pitfall	Withdrawn	The canonical Overview and gate already bind the shorthand. Repetition would reintroduce drift.

Purpose anchor contract

The largest structural defect is gone. Version 1.3 now establishes one artifact with one name and one lifecycle:

Line 71: non-routine work creates a one-sentence purpose anchor before Stage 2. Longer projects write it down; shorter work keeps it internal. Stages 3 and 4 refer back to this same anchor.

The later stages use exactly that term. The final checklist also verifies it at consequential forks. The producer and consumers now operate under the same condition.

Consequential confirmation contract

The new definition is implementable. It does not ask the skill to authenticate identity. It asks the agent to inspect the quality of the supplied context and resolve only missing, stale, or uncertain consequential facts. Independent consent requirements remain outside this purpose-check layer.

Delivery behavior

Routine work receives an internal check and quiet delivery. Visible reasoning appears only when it helps the user evaluate a material goal or judgment call. This closes the ceremony conflict without removing useful transparency.

Fable's pushback, adjudicated

Metadata placement

Fable is right. The previous recommendation to restore version, author, and platforms to the top level was wrong for Anthropic's packager.

Anthropic's current `quick_validate.py` permits these top-level keys only:

```
name, description, license, allowed-tools, metadata, compatibility
```

The validator rejects every unexpected top-level key. It deliberately ignores keys nested under metadata. The Agent Skills specification also shows author and version nested under metadata. Version 1.3 passes Anthropic's actual script unchanged.

Hermes's authoring convention favors top-level identity fields, but its documented minimum validator does not require them. That convention should inform a Hermes-specific build, not invalidate an Anthropic-compatible artifact.

Authorized-user language

Fable is right again. A Markdown skill cannot authenticate the speaker. The previous recommendation introduced a capability the artifact does not possess.

Version 1.3 uses an observable rule: explicit, current, supplied context satisfies purpose confirmation; ask only for missing, stale, or uncertain consequential facts. Authentication and consent stay with the surrounding runtime and governing domain policy.

Permissibility qualifier duplication

Fable's criticism holds. The Overview and intervention gate are the canonical sources. The pitfall's shorter "respect their call" wording remains bound by those governing rules. Repeating the qualifier everywhere would create the same drift risk the revision was trying to remove.

CORRECTION TO THE PRIOR REPORT

Recommendations 6 and 7 from the 1.2 report are withdrawn. The authorized-user wording is also withdrawn. The corrected packaging recommendation is target-specific frontmatter generation, and the corrected confirmation rule is the wording now present in 1.3.

What the review loop got right

The 89 score for 1.2 still reflected a real structural problem: later stages consumed an artifact that earlier stages did not always create. Fable preserved that finding, rejected the weak recommendations around it, and repaired the actual contract. That is the right kind of convergence.

Packaging verdicts by target

Anthropic packager

Pass. The official quick_validate.py returned Skill is valid! against the supplied 1.3 directory. Nested metadata is accepted because the script validates top-level property names and does not constrain nested keys.

Hermes

Pass at documented minimum. The file starts with valid YAML frontmatter, contains name and description, stays below the description and file-size limits, and has a non-empty body. Hermes can load the nested metadata mapping.

Convention note. Hermes peers commonly place version, author, and platforms at the top level. That is a local style preference, while Anthropic’s packager makes the opposite shape mandatory. A universal source file cannot satisfy both conventions literally.

Agent Skills reference validator

Fail for the supplied artifact. The official skills-ref implementation reported invalid frontmatter at platforms: [linux, macos, windows] because it rejects flow-style list syntax in metadata. After converting platforms to a string, it rejected the flow-style tags list. A flattened string-only metadata variant passed.

It also requires the parent directory name to equal the skill name. A temp build under purpose-check/ passed; the same content under purpose-check-1.3/ failed because the directory did not match name: purpose-check.

Correct build strategy

Keep one behavioral body and render target-specific frontmatter at package time:

Target	Frontmatter	Package directory
Anthropic	Keep identity fields under metadata. Current shape passes.	purpose-check-1.3/ accepted by current packager, though purpose-check/ is cleaner.
Hermes	Either current nested shape or a Hermes-specific top-level identity variant.	purpose-check/ preferred.
Strict Agent Skills	Use only allowed top-level fields. Flatten metadata values to strings.	Must be purpose-check/.

A ten-line packaging script or template is enough. The behavioral body should remain identical across builds.

Behavioral quality review

Trigger and routing

The description remains 386 characters and names the conditions that justify activation: material ambiguity, goal mismatch, consequential forks, and consequential domains. Routine work is explicitly excluded from visible ceremony.

The gate is ordered by consequence and maps cleanly to five behaviors: proceed, state an assumption, ask one useful question, flag a mismatch, or confirm consequential facts. Highest-consequence precedence prevents competing rows from producing mixed behavior.

Stage coherence

The sequence now holds together:

1. Stage 1 identifies the audience and purpose, then creates the purpose anchor for non-routine work.
2. Stage 2 defines a verifiable end state.
3. Stage 3 compares material forks against the anchor.
4. Stage 4 returns to the anchor and checks for drift.
5. Verification confirms that material forks were handled against the same anchor.

Every later reference has a producer. No stage depends on a process file for short work.

Safety and autonomy

The skill maintains one canonical rule against moral substitution in the Overview. The mismatch gate respects informed user choice only when the request remains permissible. Consequential confirmation supplements independent safety, privacy, legal, authorization, and consent checks.

This division is sound. Purpose alignment does not pretend to be authentication, permission enforcement, or legal judgment.

Writing quality

The prose has a distinct voice and avoids generic requirements language. Definitions arrive before use. Examples are specific. The Greek vocabulary gives the model memorable internal handles while the skill explicitly keeps those terms out of user-facing language.

The longest sentence is the consequential-domain definition on line 27. It carries confirmation semantics, context reuse, default refusal, and independent-consent boundaries in one paragraph. It is dense, but co-locating the governing rule is better than redistributing qualifiers across the document.

EDITORIAL VERDICT

Do not split line 27 merely to make it prettier. If edited, preserve one canonical location for the complete consequential-domain contract.

Validation and scenario analysis

Mechanical checks

Check	Result	Details
Archive integrity	Pass	One unencrypted Markdown payload; no symlinks or unsafe paths.
Anthropic quick validator	Pass	Skill is valid! on the supplied directory.
Hermes documented minimum	Pass	Valid YAML, required fields, description limit, non-empty body.
Structural contract verifier	Pass	Purpose anchor produced and consumed consistently; old Stage 1 block term absent.
Agent Skills reference validator	Fail	Metadata collection shapes rejected; versioned directory name also mismatches skill name.
Flattened reference build	Pass	String-only metadata under purpose-check/ validated successfully.

Scenario matrix

Scenario	Rule-consistent behavior
Clear reversible task	Brief internal check, no question, quiet delivery.
Minor safe ambiguity	State one-line assumption and proceed.
Dashboard missing audience	Ask one high-value question with recommendation or options.
Consequential task with complete current context	Treat supplied facts as confirmation, subject to independent consent rules.
Consequential task with stale authority detail	Ask only for the stale or uncertain fact.
Medium task with a material fork	Create an internal purpose anchor, compare the fork, and surface the user-reaching difference.
Likely goal mismatch	Flag once, offer the better fit, and respect a permissible informed choice.

No written-rule conflicts remain in these scenarios.

Final proof still missing

The skill has passed static and deterministic review. It has not yet been benchmarked through live model runs. The last five points should come from evidence rather than more prose edits.

Use at least 24 prompts across routine work, safe defaults, deliverable-changing ambiguity, consequential work with complete context, consequential work with missing context, and mid-task forks. Measure unnecessary questions, missed clarifications, redundant confirmations, visible ceremony, and purpose-anchor retention.

Final assessment

BOTTOM LINE

Version 1.3 is behaviorally ready. It resolves the core structural issue, uses implementable confirmation semantics, and preserves one canonical safety boundary. Ship it for Anthropic and Hermes. Add target-specific packaging before claiming strict Agent Skills portability.

The fresh score is **95 / 100**. The six-point increase from 1.2 comes from closing actual contracts rather than adding sections. The remaining deduction is split between packaging portability and the lack of live behavioral benchmark results.

Fable's response improved the review itself. It correctly separated a local Hermes convention from Anthropic's enforced schema, rejected a capability the skill cannot possess, and defended canonical-rule discipline against qualifier duplication. Those corrections are now incorporated into this report.

Review method and confidence

The review used the actual ZIP payload, a line-level comparison against 1.2, authoritative Agent Skills documentation, Anthropic's live validator source, Anthropic quick-validator execution, Hermes minimum-validator reproduction, the official `skills-ref` implementation under Python 3.11, temporary build variants, archive safety checks, a structural contract script, and a seven-scenario rule matrix.

- **Confirmed:** all four behavioral findings from 1.2 are resolved.
- **Confirmed:** Anthropic's validator passes the supplied artifact.
- **Confirmed:** strict `skills-ref` rejects the supplied metadata shapes and versioned directory.
- **Confirmed:** a flattened string-only build under `purpose-check/` passes `skills-ref`.
- **High confidence:** the written behavior is internally coherent.
- **Medium confidence:** live model behavior will match the rules. A benchmark remains necessary.

Artifact identity

Attachment SHA-256: `ddbda5be3fb1b26ae260413761db52fabd0f9b3543bbc2bb0aaa8deaad8753b`

SKILL.md SHA-256: `e9be4ca2127959a55ae98316802fe776a2e8207adf6a2ce1702ac1d08abb1cb4`

Final verdict: approve behavior, add target-specific packaging.